

Mutian He

Data Engineer Intern

Email: mhe3@scu.edu | Phone: 408-548-1060 | San Jose, CA

LinkedIn: www.linkedin.com/in/mutian-he

Github: github.com/austin10231

Portfolio: www.mutianhe.com

SUMMARY

Data Engineering graduate student with experience building real-time data pipelines and scalable cloud data lakes using Python, SQL, Kafka, and AWS. Skilled in designing ETL workflows and integrating ML pipelines for data-driven applications.

EDUCATION

Santa Clara University, Leavey School of Business

Santa Clara, US

Master of Science in Information Systems, GPA 3.59

Dec 2026

- Data Analytics with Python, Cloud Computing, Database Systems, Machine Learning, Deep Learning, NLP

University of Glasgow, School of Computing Science

Glasgow, UK

Bachelor of Science in Computer Science (BSc CS), GPA 3.66

Jun 2024

- Data Structures & Algorithms, Database Systems, Software Engineering, Human-Computer Interaction

EXPERIENCE

Data Science Intern

ViiaRaa Inc. — Palo Alto, CA — Mar 2026 – Present

- Analyze time-series behavioral datasets to identify patterns and trends supporting personalized wellness insights.
- Develop predictive machine learning models and statistical analyses to support product decision-making.
- Collaborate with cross-functional teams to design experiments and evaluate data-driven product improvements.
- Build data visualizations and dashboards to communicate insights and support product strategy.

PROJECTS

Cloud-Based Stock Data Pipeline –Data Engineering (2026) [Github](#)

- Designed and deployed a real-time streaming data pipeline using **Apache Kafka** on **AWS EC2** for stock market data ingestion.
- Architected a **scalable AWS data lake** using Amazon S3 and AWS Glue Data Catalog for schema management.
- Implemented **automated ingestion pipelines** to stream stock market data into an **AWS-based data lake architecture**.
- Enabled serverless **SQL analytics** on streaming datasets using **Amazon Athena** over data stored in **S3**.

Risk Change Alert System – Cloud & AI Engineering (2026) [Live Demo](#) | [Github](#)

- Built an **end-to-end document processing pipeline** to ingest and parse SEC 10-K filings (HTML + PDF) using **BeautifulSoup** and **AWS Textract**.
- Integrated **AWS Bedrock (Nova Lite)** to **auto-classify** risks into 13 categories and generate executive-level summaries, enabling AI-powered year-over-year and cross-company risk comparison.
- Deployed a **full-stack Streamlit application on Streamlit Cloud with AWS S3** for persistent storage and multi-dimensional filtering across company, year, and industry.

Instacart Reorder Prediction Pipeline – Machine Learning (2026) [Github](#)

- Developed a modular ML data pipeline including **ETL, feature engineering, model training**, and inference driven by YAML configuration.
- **Eliminated data leakage** via strict temporal splitting, ensuring robust evaluation under real-world conditions.
- **Deployed an inference script** (predict.py) to load serialized models (.joblib) and simulate real-time predictions.

SKILLS

Programming: Python, SQL (MySQL, SQLite, PostgreSQL, MongoDB), Java

Data Engineering: Apache Kafka, Apache Airflow, ETL/ELT Pipelines, Streaming Architecture, Data Modeling

Machine Learning: Scikit-learn, Feature Engineering, Cross-Validation (K-Fold), Hyperparameter Tuning, Model Evaluation (ROC-AUC, F1, Recall), Random Forest, PCA, K-Means

Cloud & Platforms: AWS (EC2, S3, Glue, Athena, RDS, Lambda), Data Lake Architecture

Tools: Git, Docker, Linux, Jenkins, Streamlit, Flask, VS Code, Jupyter Notebook